

# Introduction to AWS EC2 F1

SDx 2018.2



# Objectives

- > **After completing this module, you will be able to:**
  - >> Describe AWS EC2 services
  - >> List some of the benefits of using AWS EC2 F1 instance
  - >> State application areas where acceleration is beneficial
  - >> Describe SDAccel development flow on AWS EC2 F1
  - >> Explain what are AMI and AFI
  - >> Create an account on AWS (Appendix)
  - >> Create and access an instance (Appendix)


# Outline

- > Amazon AWS EC2 F1
- > SDAccel Development Flow on EC2 F1
- > Accessing an AWS EC2 Instance
- > Connecting to an Instance
- > Summary
- > Lab1 Intro
- > Appendix



# Introducing Amazon AWS EC2 F1



- > Amazon AWS EC2 F1 is a compute instance with Xilinx FPGAs which can be programmed to create custom hardware accelerated applications
- > AWS EC2 F1 instances are easy to program and come with everything needed to develop, simulate, debug, and compile hardware accelerators
- > Once a FPGA design is complete, it can be registered as an Amazon FPGA Image (AFI), and deployed to F1 instance in just a few clicks



Host CPU

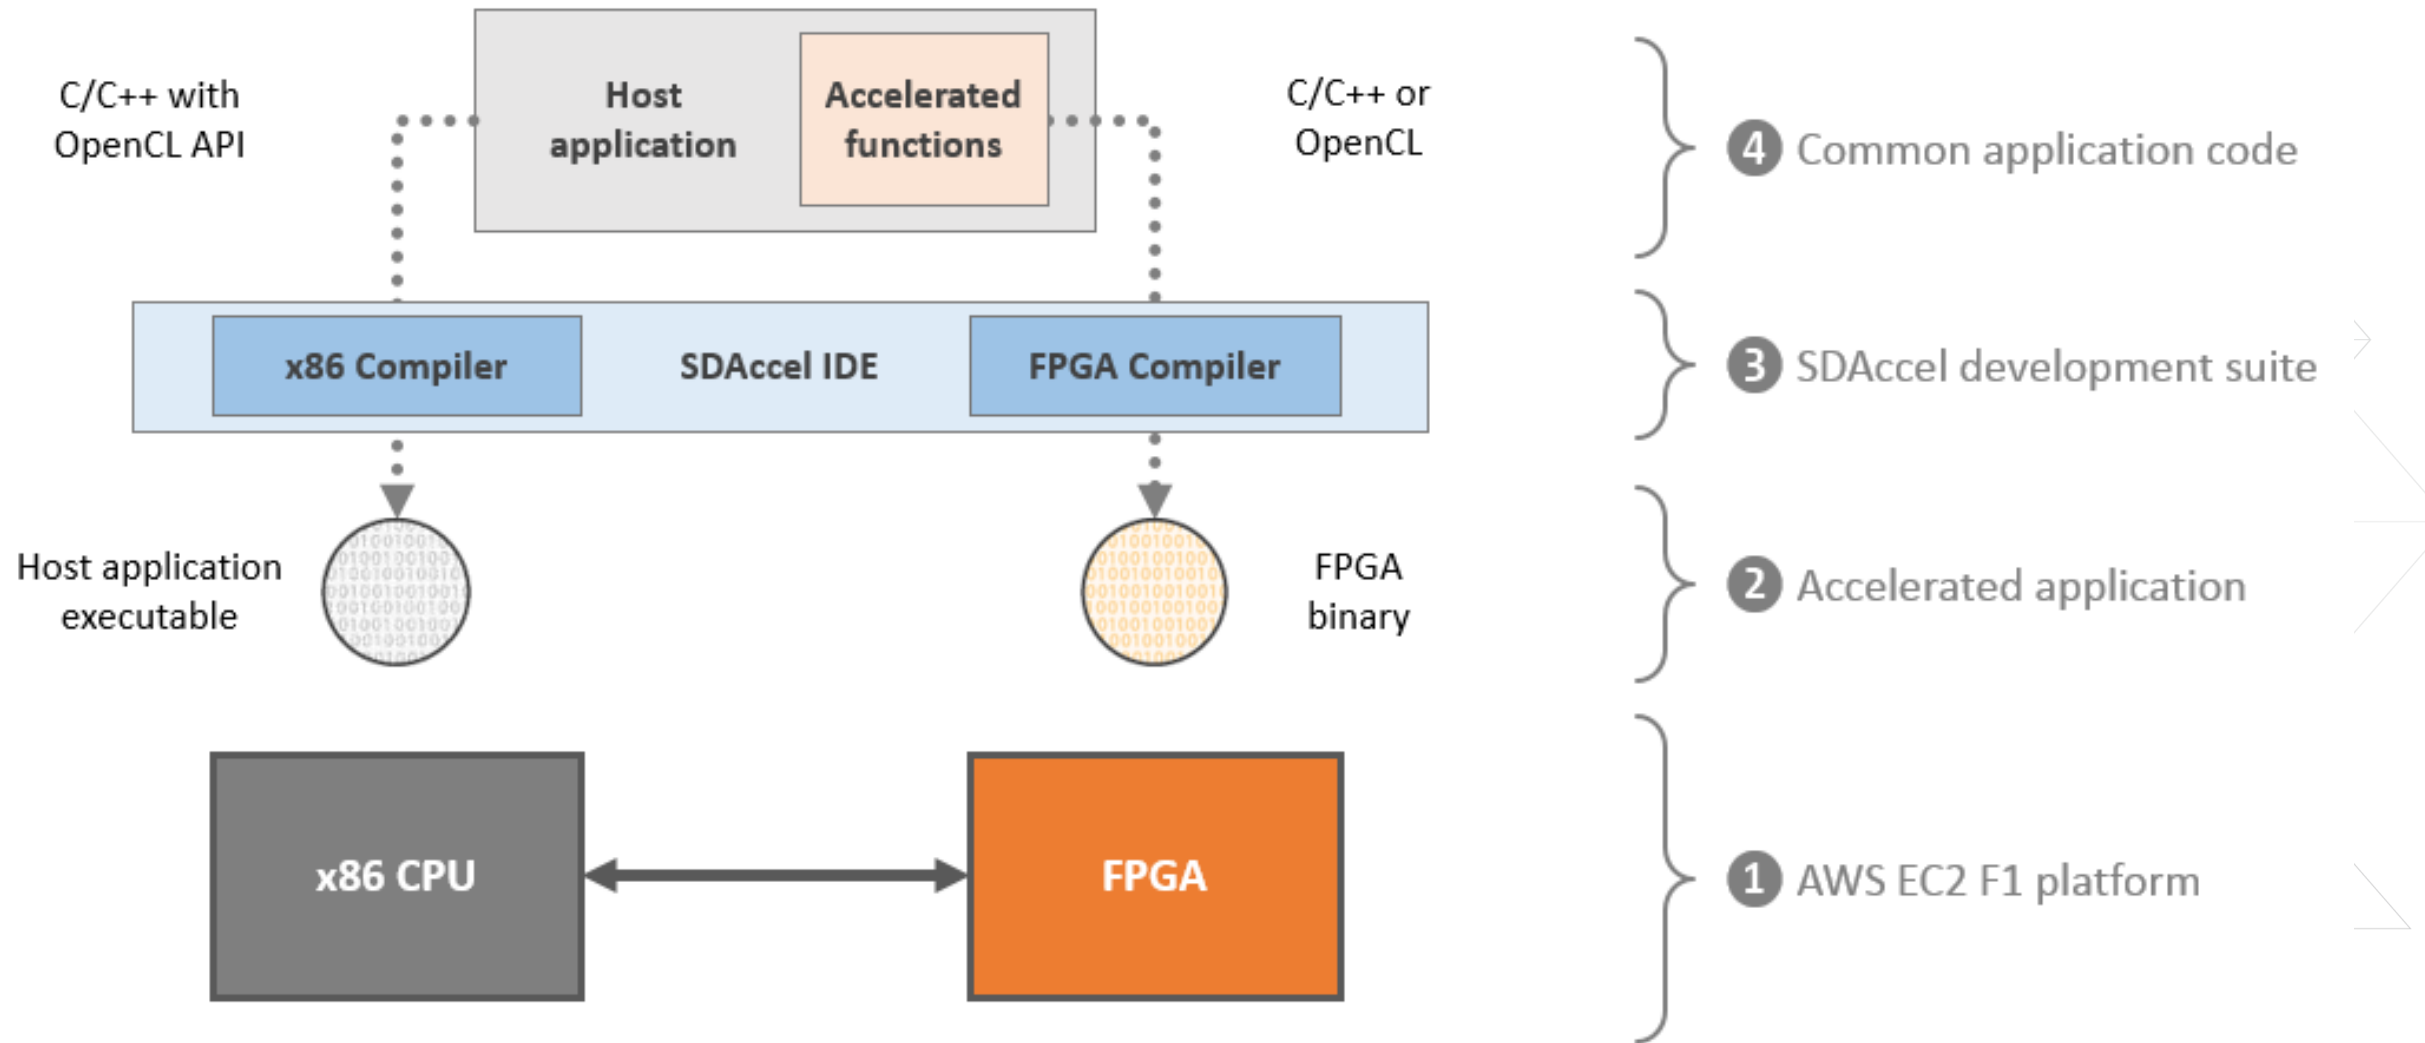


Xilinx FPGA

- > **Up to 8 Xilinx UltraScale+ 16nm VU9P FPGA devices in a single instance**
  - >> f1.16xlarge size provides:
    - 8 FPGAs, each with over 2 million customer-accessible FPGA programmable logic cells and over 5000 programmable DSP blocks
  - >> Each of the 8 FPGAs has 4 DDR-4 interfaces, with each interface accessing a 16GiB, 72-bit wide, ECC-protected memory

Instance Size	FPGAs	DDR-4 (GiB)	vCPUs	Instance Memory (GiB)	NVMe Instance Storage (GB)	Network Bandwidth
f1.2xlarge	1	4 x 16	8	122	1 x 470	Up to 10 Gbps
f1.16xlarge	8	32 x 16	64	976	4 x 940	25 Gbps

# AWS EC2 F1 Instance SDAccel Flow



# Workloads Dependent Acceleration



Video

10x




Data analytics

90x



Genomics

100x



Machine Learning

40x



& more



Reconfigure.io



# Terminology

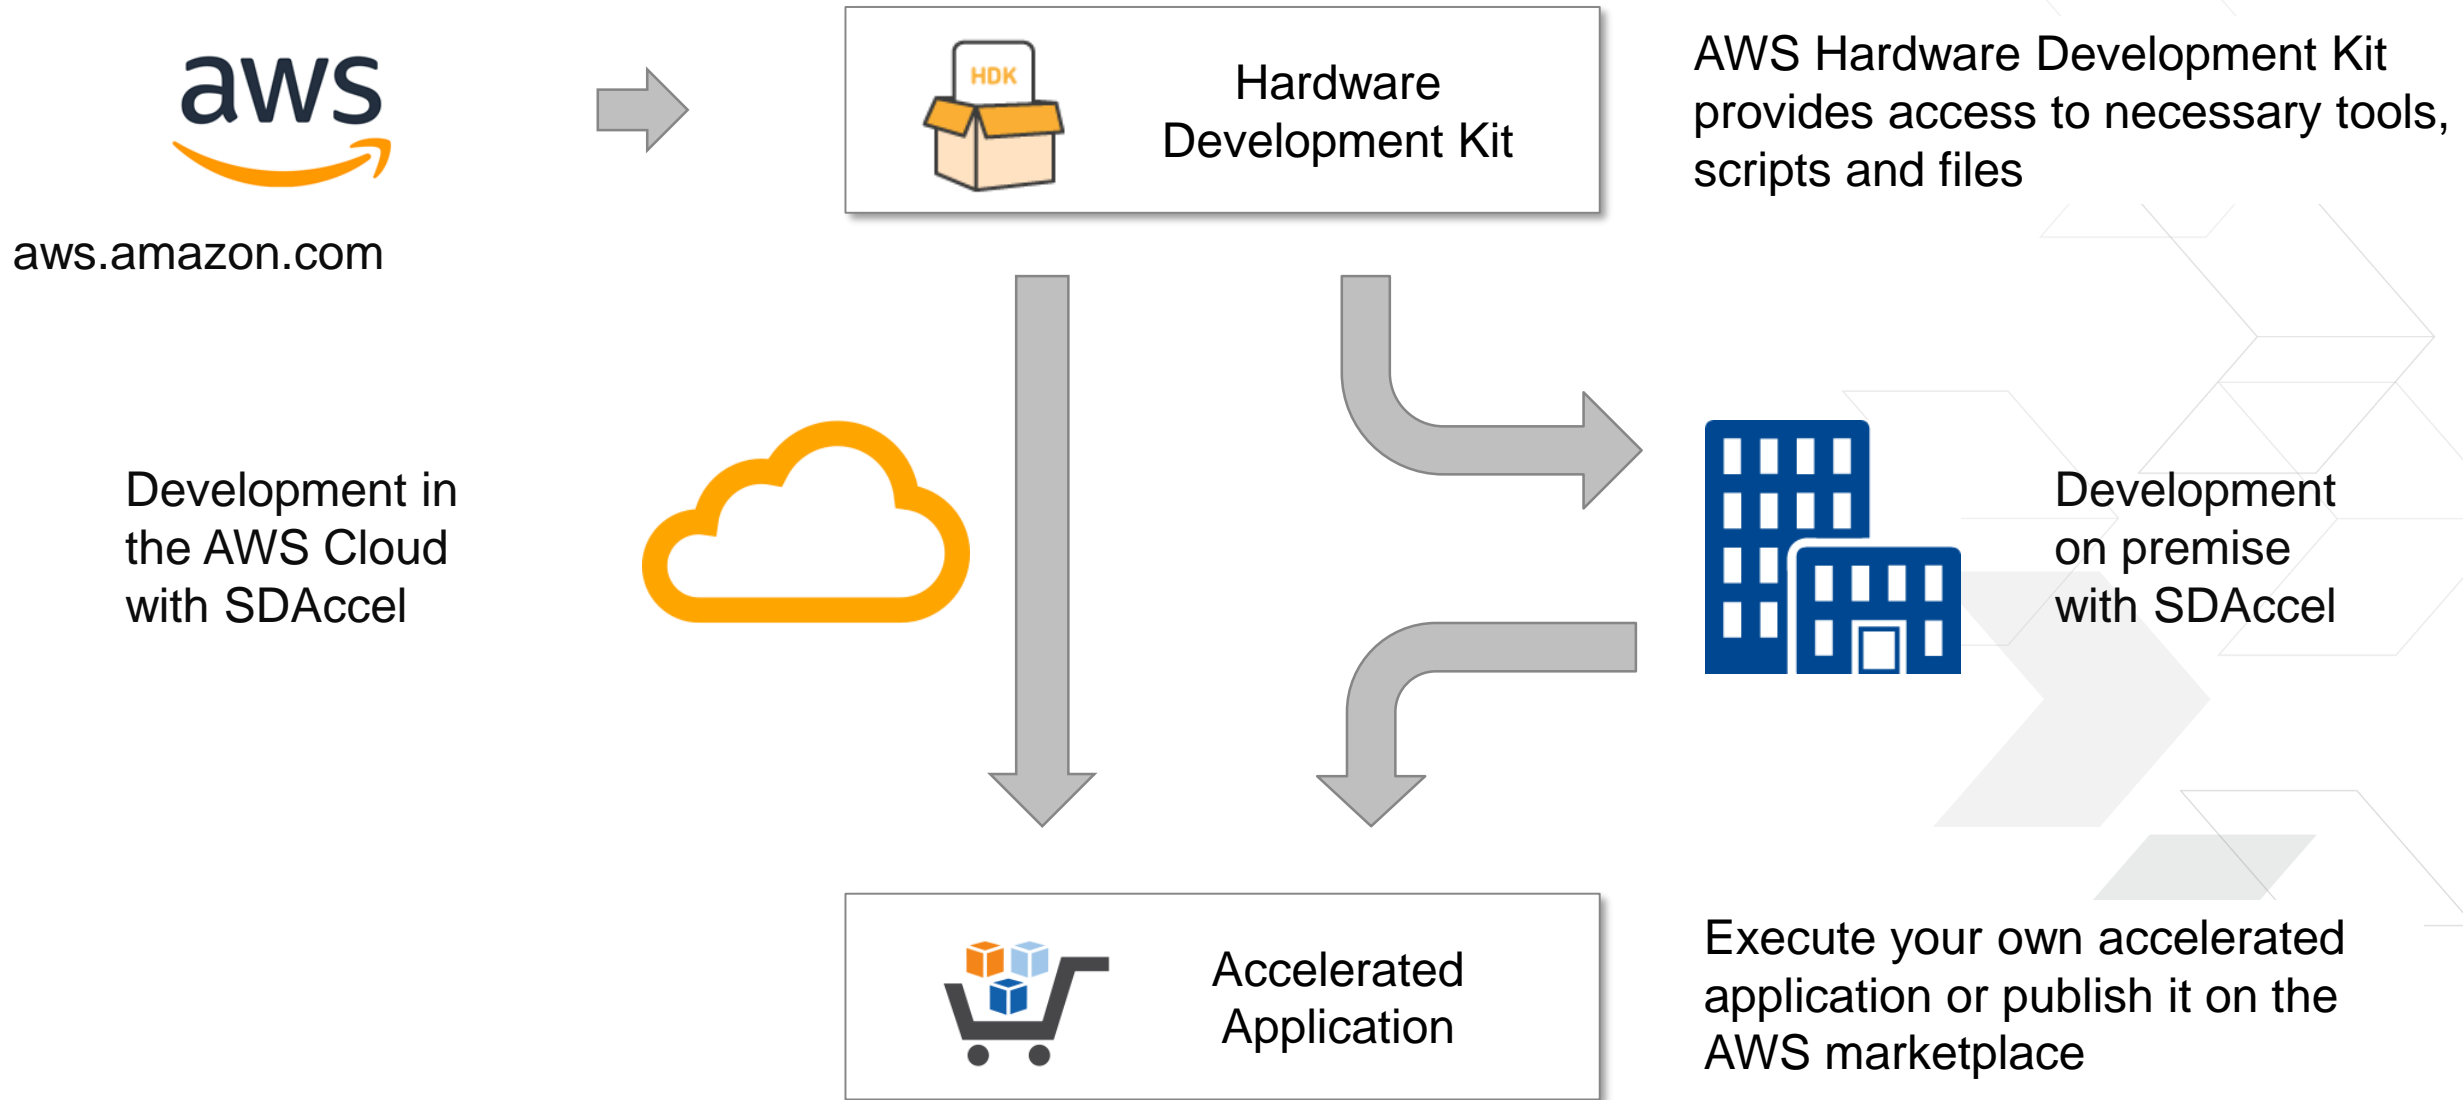
- > **AMI: An Amazon Machine Image is a template that contains a software configuration (for example, an operating system, an application server, and applications)**
  - >> Provides the information required to launch an instance, which is a virtual server in the cloud
  - >> Must specify a source AMI when you create an instance
  - >> Multiple instances can be launched using the same source AMI
- > **AFI: An Amazon FPGA Image is the compiled registered design, securely stored**
  - >> Secured, encrypted and dynamically loaded in the FPGA - can't be copied or downloaded
  - >> Can be associated with an AMI and offered on the AWS Marketplace




# SDAccel Development Flow on EC2 F1



# Amazon F1 Development Flow

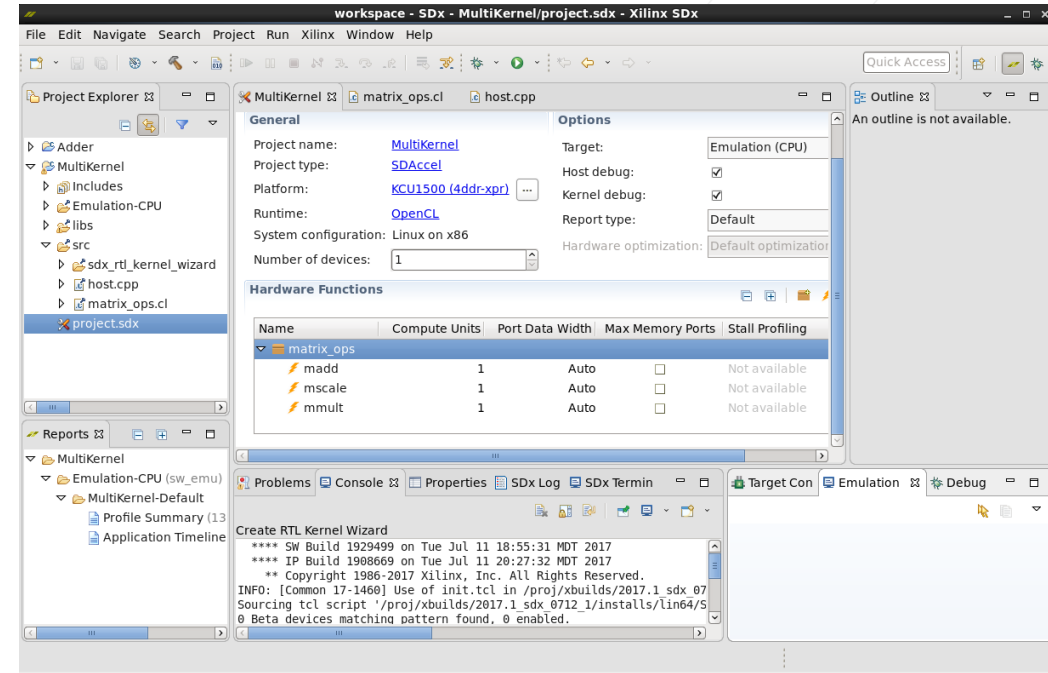


# The SDAccel Development Environment

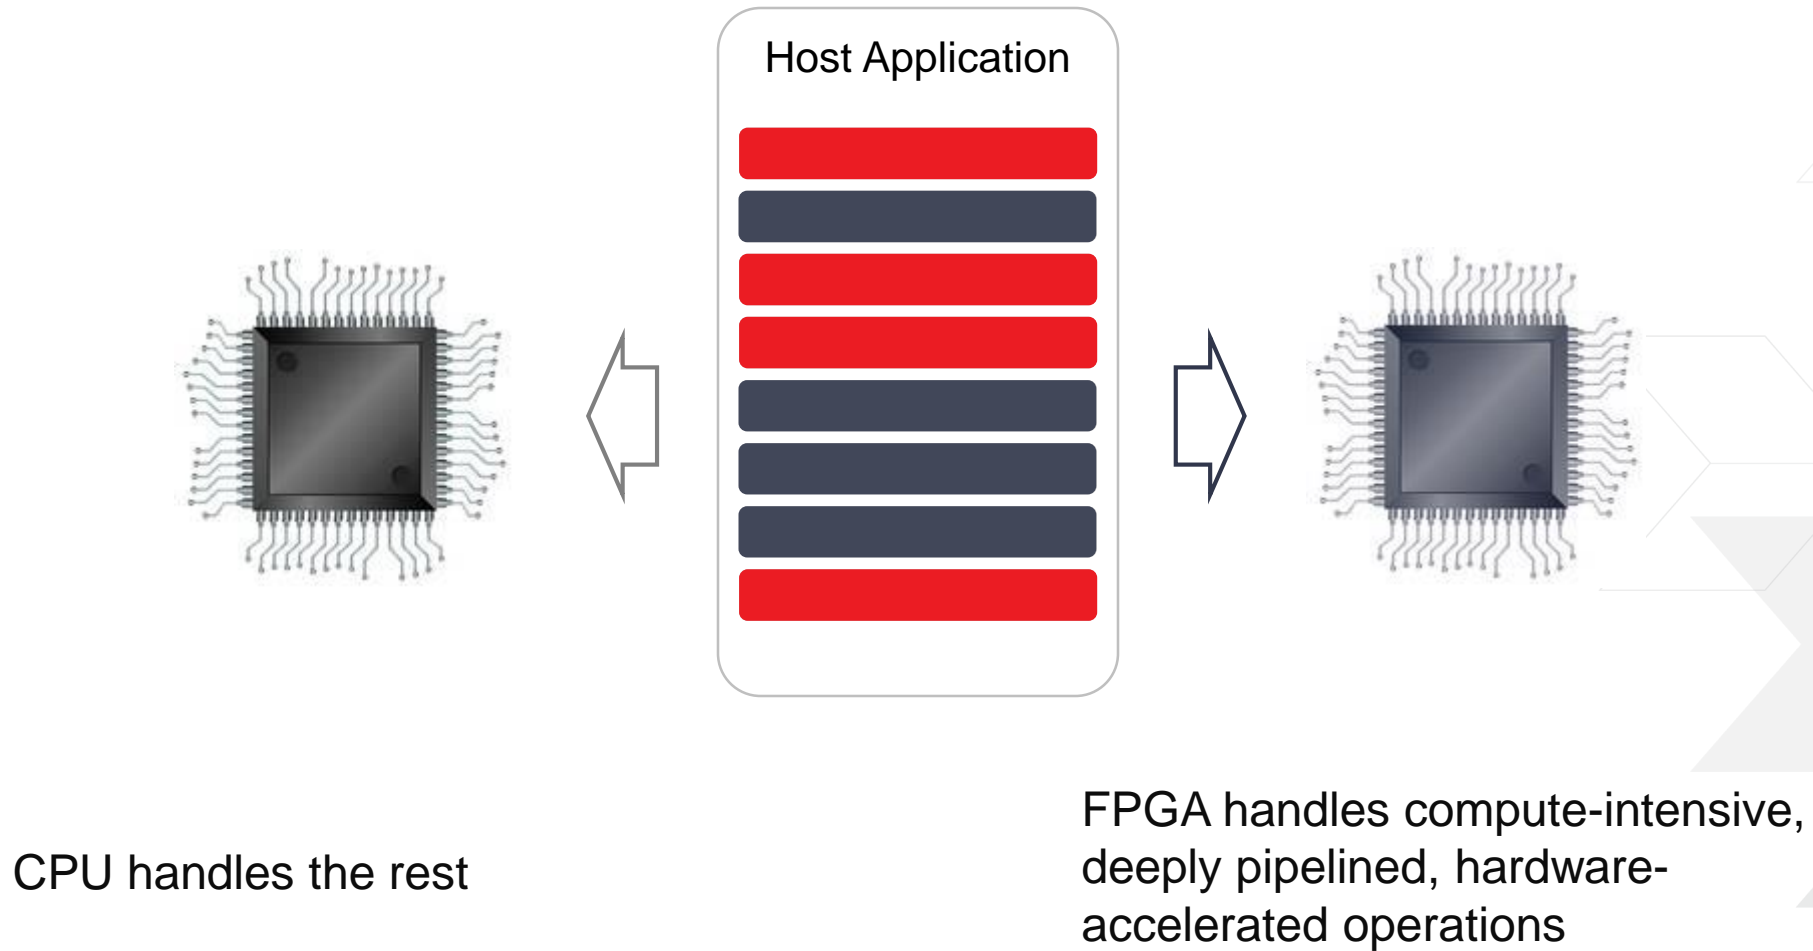


# The SDAccel Development Environment

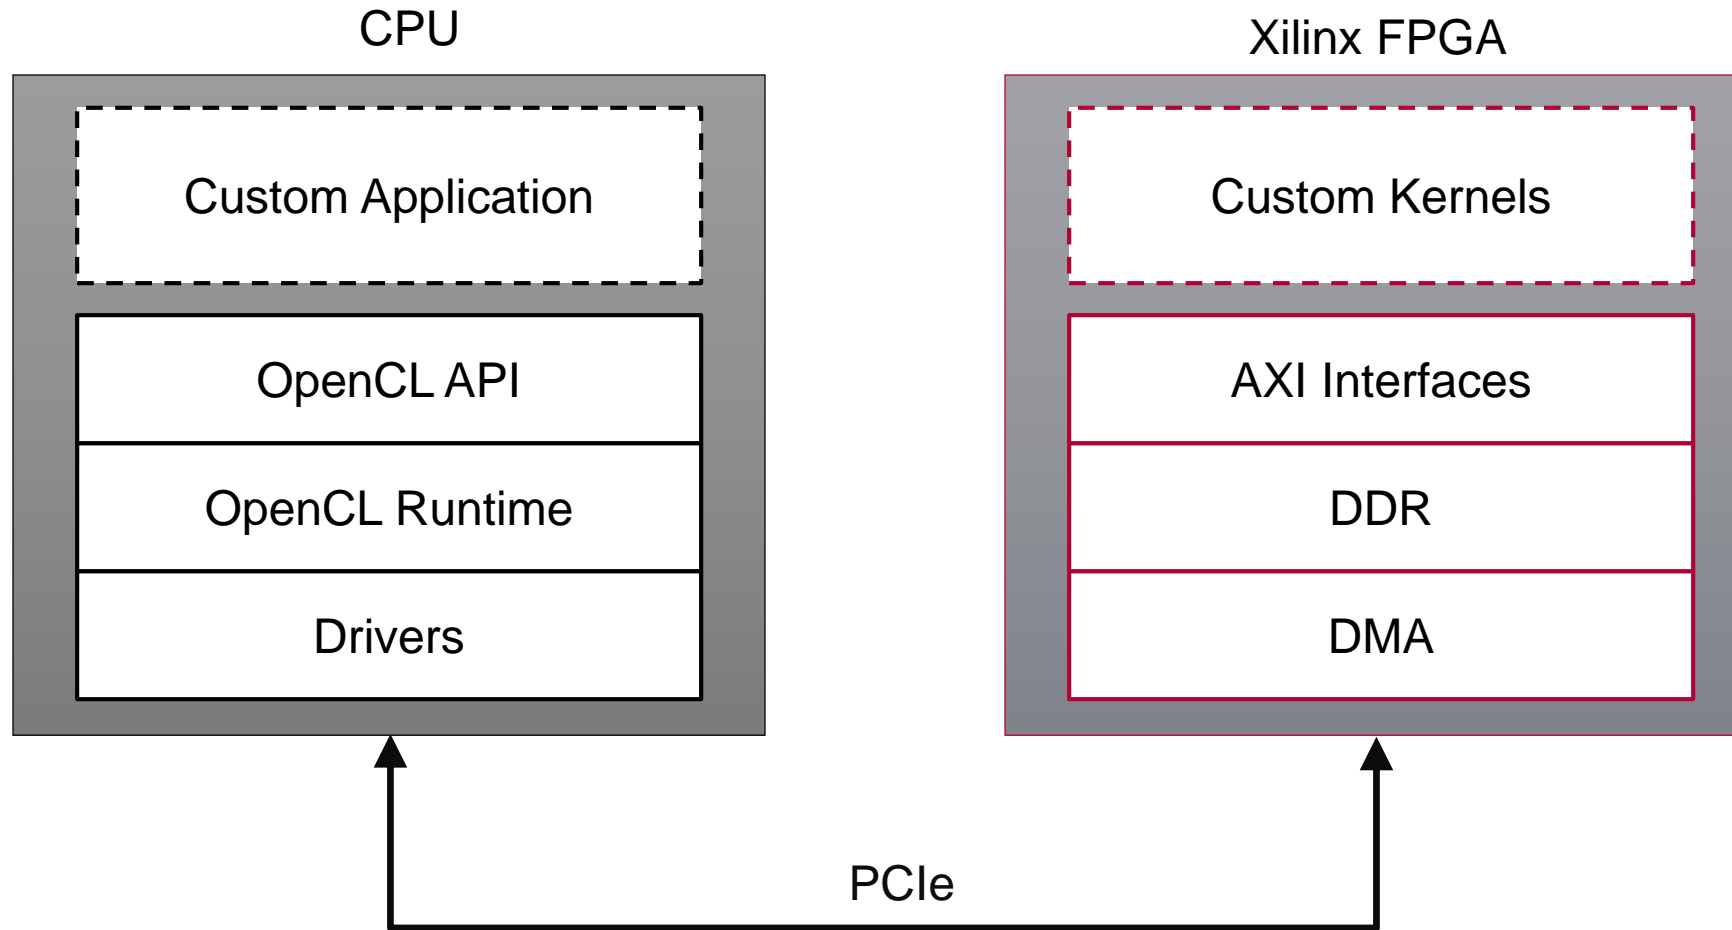
- > Fully integrated Eclipse-based environment
- > Develop, profile and deploy applications accelerated with Xilinx FPGAs
- > Concurrent programming of the host application and FPGA kernels
- > Automatic hardware execution flows
- > Built-in debug, profiling and performance analysis tools



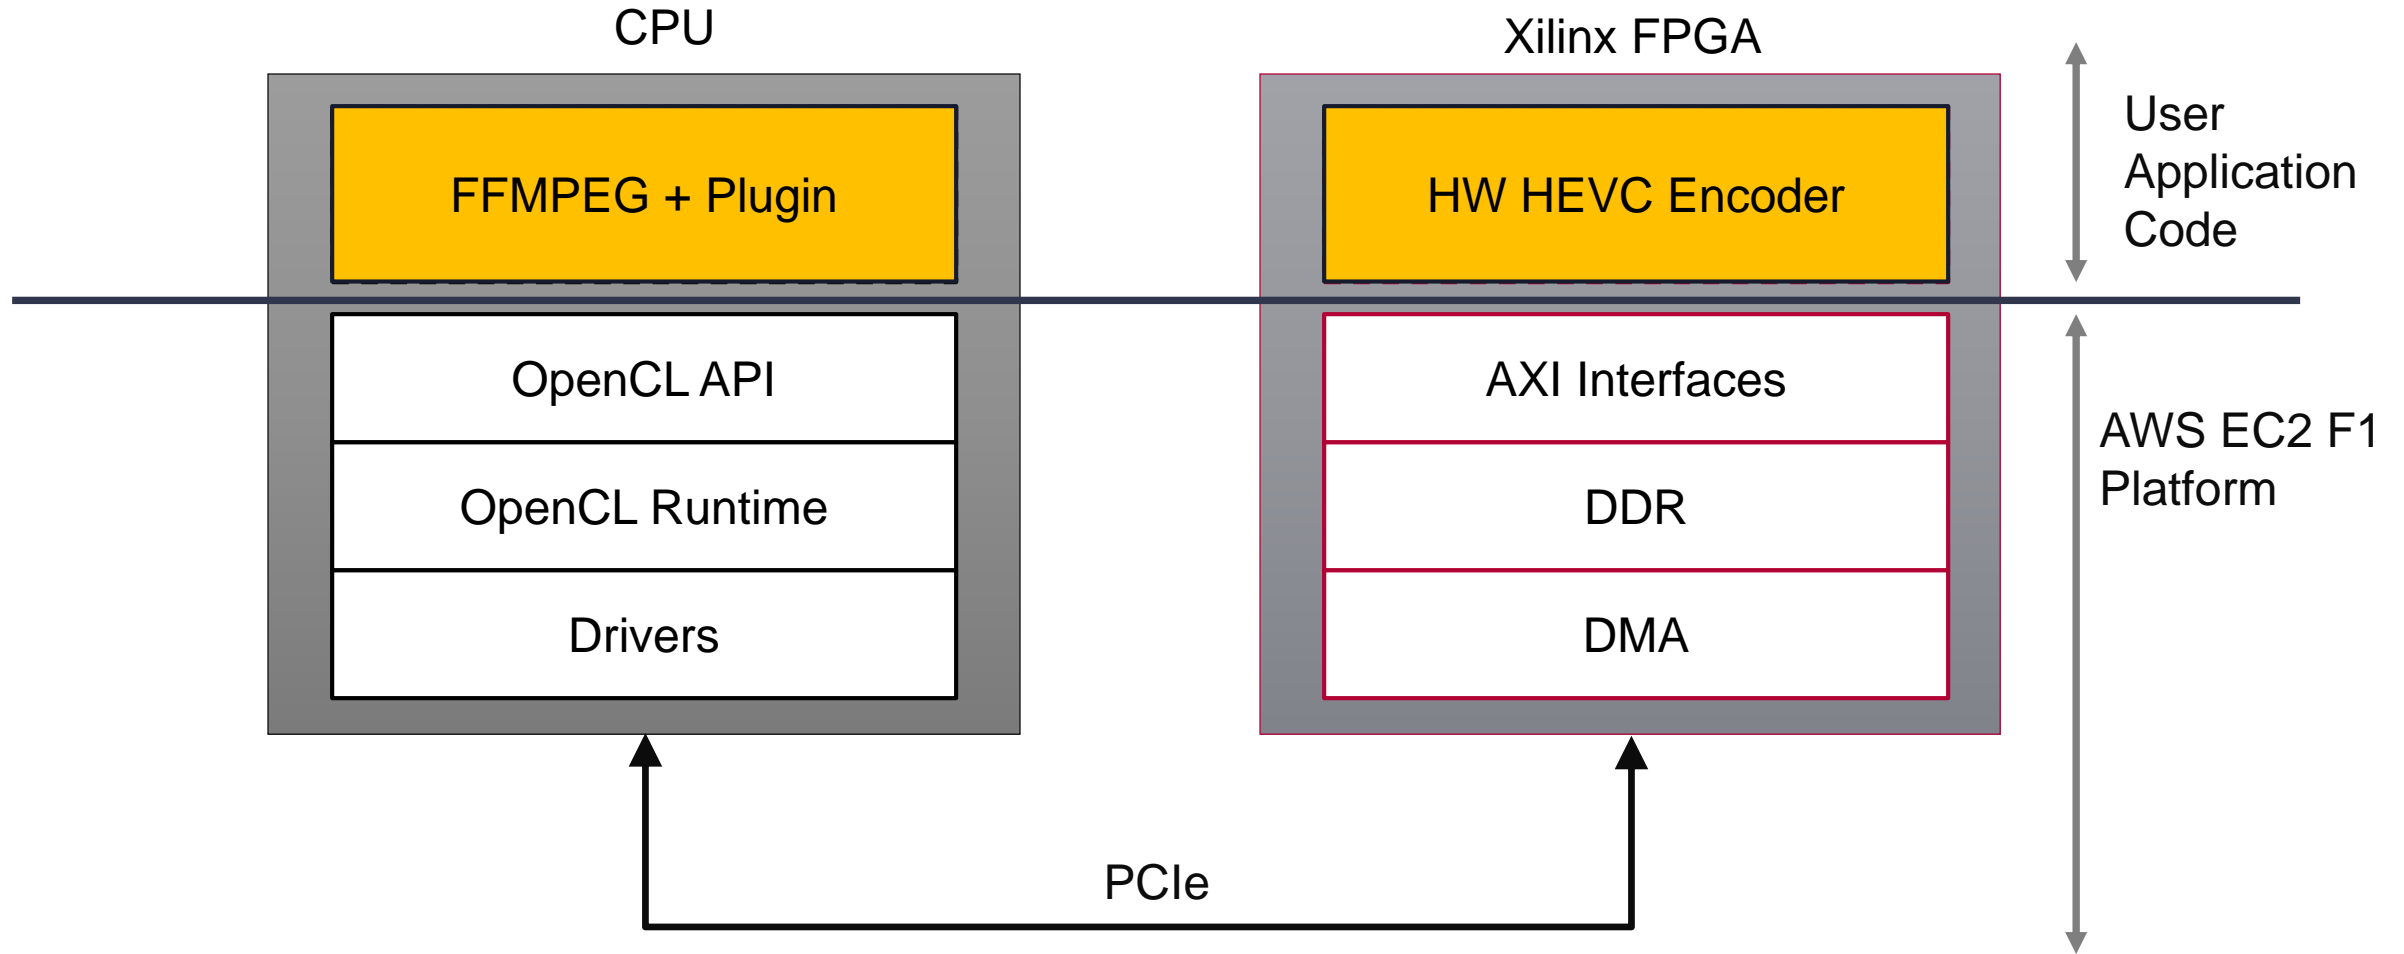
# How FPGA Acceleration Works on AWS



# AWS EC2 F1 Platform Model



# Example of a Custom Video Encoding Application

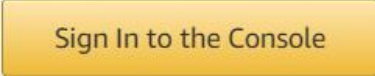


# Accessing an AWS EC2 Instance





# Login Procedure using Provided Instance

- > **Click on the provided link to your preconfigured instance**
  - >> Otherwise go to **aws.amazon.com**, click on 
- > **Use the provided Account ID**
  - >> `xilinx-aws-f1-developer-labs`
- > **Enter provided user name**
  - >> `usern`
- > **Enter the following password**
  - >> *Given by instructor*
  - >> *For centos user on RDP: Given by instructor*
- > ***Precompiled FPGA binaries***



Account ID or alias

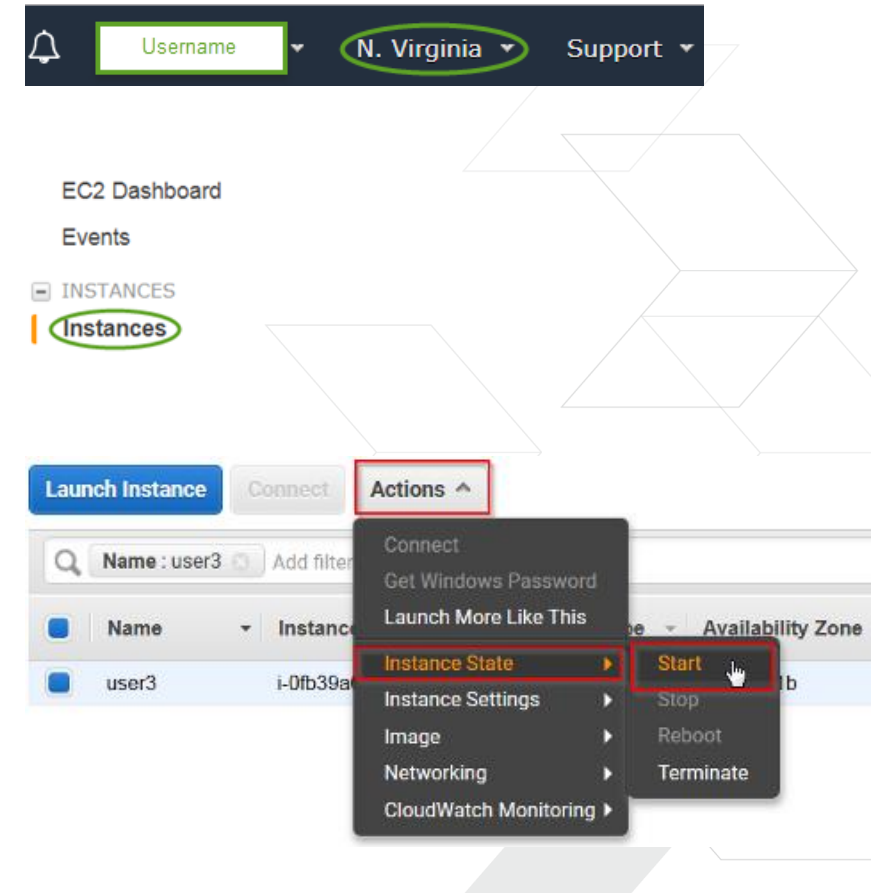
IAM user name

Password

Sign In

# Starting an Instance on EC2 F1

- > Make sure that you have selected *N. Virginia* as a region
- > Click on the instances in the left pane
- > Select your instance and then start it
  - >> Since the AMI used already has an instance available, filter your instance by typing your username in the Search field under the Launch Instance button



# Connecting to an Instance



# Connection setup

- > You can communicate with the instance using
  - >> Command line through PuTTY or Git Bash
  - >> GUI through remote desktop connection
- > Select the running instance and you will see the console window with various information

Instance: **i-0ea02fdf9d5fd32b1** Public DNS: **ec2-54-209-156-37.compute-1.amazonaws.com**

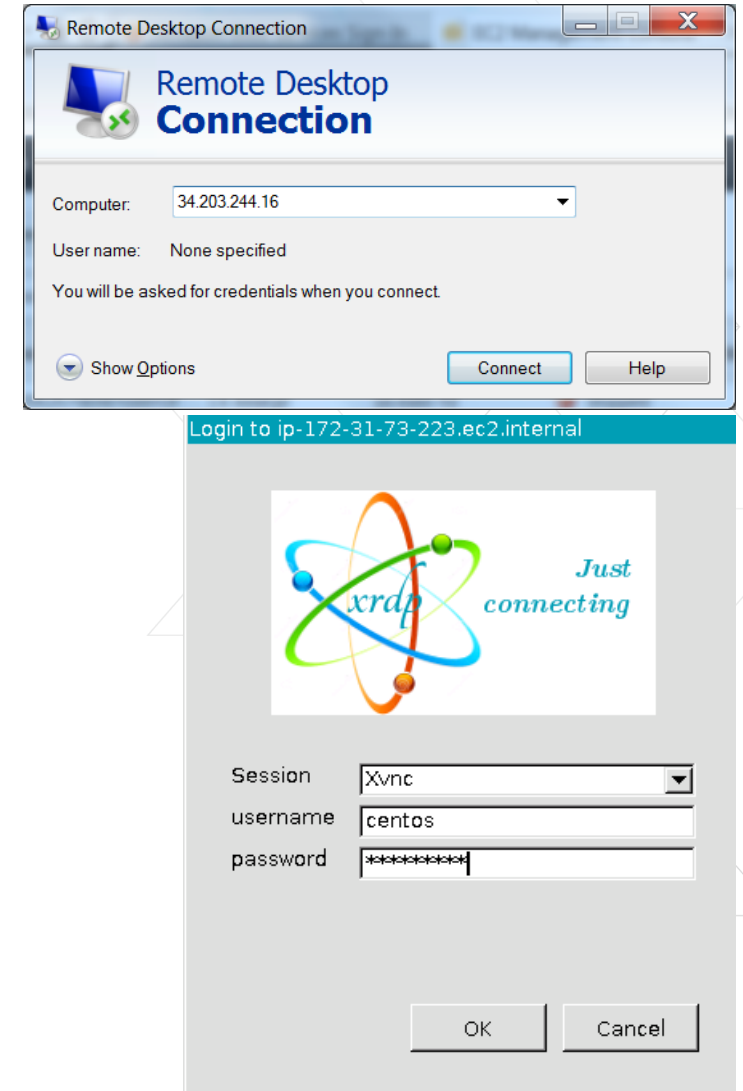
Description	Status Checks	Monitoring	Tags	Usage Instructions
Instance ID	i-0ea02fdf9d5fd32b1			
Instance state	running			
Instance type	f1.2xlarge			
Elastic IPs				
Availability zone	us-east-1c			
Security groups	FPGA Developer AMI-1-3-3-AutogenByAWSMP-3 <a href="#">view inbound rules</a>			
Scheduled events	No scheduled events			
AMI ID	GUI_test (ami-b169cdcb)			
Platform	-			
IAM role	-			
Key pair name	iccd_users			
Owner	331708262607			
Launch time	October 30, 2017 at 10:09:51 AM UTC-7 (less than one hour)			
Public DNS (IPv4)	ec2-54-209-156-37.compute-1.amazonaws.com			
IPv4 Public IP	54.209.156.37			
IPv6 IPs	-			
Private DNS	ip-172-31-70-48.ec2.internal			
Private IPs	172.31.70.48			
Secondary private IPs				
VPC ID	vpc-c9ba6bb0			
Subnet ID	subnet-0478a028			
Network interfaces	eth0			
Source/dest. check	True			
EBS-optimized	True			
Root device type	ebs			

used in ssh connection with PuTTY or Git Bash  
as well as for remote desktop connection

used for remote desktop connection

# Starting a Remote Desktop Session

- > Start *Remote Desktop Connection* program
- > Enter *IPv4 Public IP* address available in the console
- > Click **Connect**
- > Click **Yes**
- > In the desktop session
  - >> Enter **centos** in the *username* field
  - >> Enter <user password> as you had set that in PuTTY session
  - >> Click **OK** to run the session
- > You can now execute **GUI/Windows based tasks**



# Remote Desktop Client

- > **Before connecting, set your remote desktop client to use 24-bit for color depth**
  - >> Windows client: In the bottom-left corner of connection prompt, click Options, select the Display tab and set Colors to True Colors (24 bit)
- > **On Windows: press the Windows key and type "remote desktop"**
  - >> You should see the "Remote Desktop Connection" in the list of programs
  - >> Alternatively you can also simply type mstsc.exe in the Windows run prompt
- > **On Linux: RDP clients such a Remmina or Vinagre are suitable**
- > **On MacOS: use the Microsoft Remote Desktop v8.0.43 (that version offers color depth settings) from the Mac App Store**

# Summary



# Summary

- > **AWS EC2 F1 is an elastic cloud computing instance that uses Xilinx FPGAs for custom hardware accelerated applications**
- > **AWS EC2 F1 development flow options**
  - >> Develop in the AWS cloud with the SDAccel development environment
  - >> Develop on-premise in the SDAccel environment and execute the application on AWS
- > **Accelerated applications can easily be put on the AWS Marketplace, offering a rapid path to monetization**



# Lab Intro



# Lab Intro

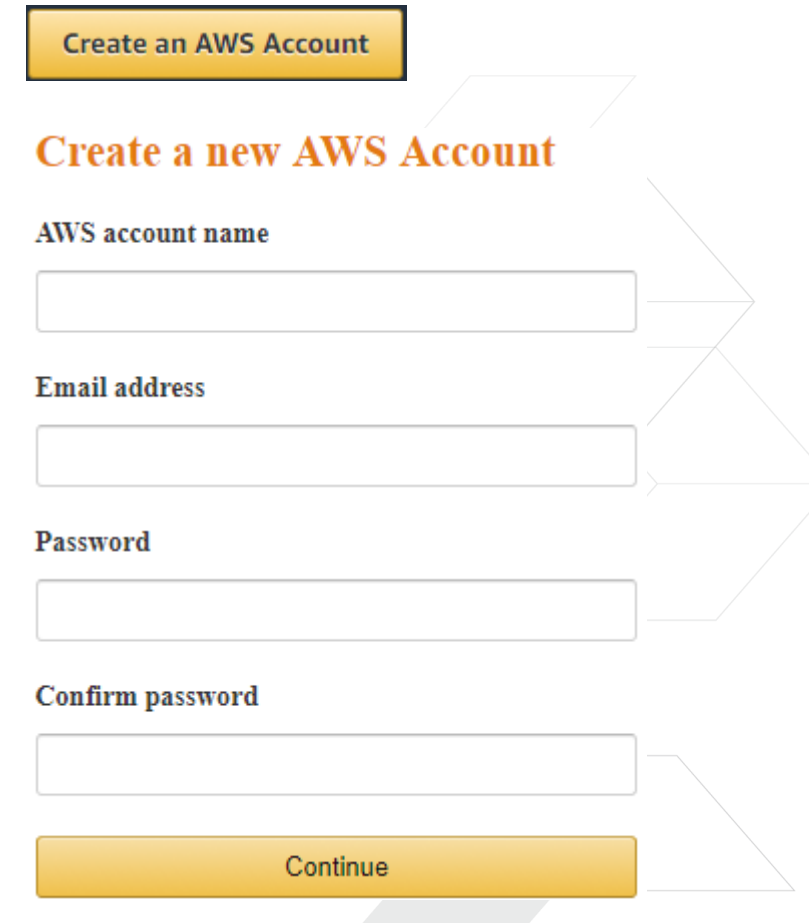
- > In this lab you will start an EC2 F1 instance using the preconfigured FPGA Developer AML and connect to it using a remote desktop client. Once connected, you will open a terminal window on the RDP client and source the environment variables.

# Appendix: Account Creation on AWS



# Creating an AWS Account

- > Go to <https://aws.amazon.com> and then click on
  - >> Free account is really for non-FPGA services
- > Enter your choice of account name, email address password and click on Continue



The image shows a screenshot of the AWS account creation page. At the top, there is a yellow button labeled "Create an AWS Account". Below it, the heading "Create a new AWS Account" is displayed in orange. The form consists of four input fields: "AWS account name", "Email address", "Password", and "Confirm password". Each field is followed by a yellow button labeled "Continue". The form is decorated with light gray geometric shapes in the background.

Create an AWS Account

Create a new AWS Account

AWS account name

Email address

Password

Confirm password

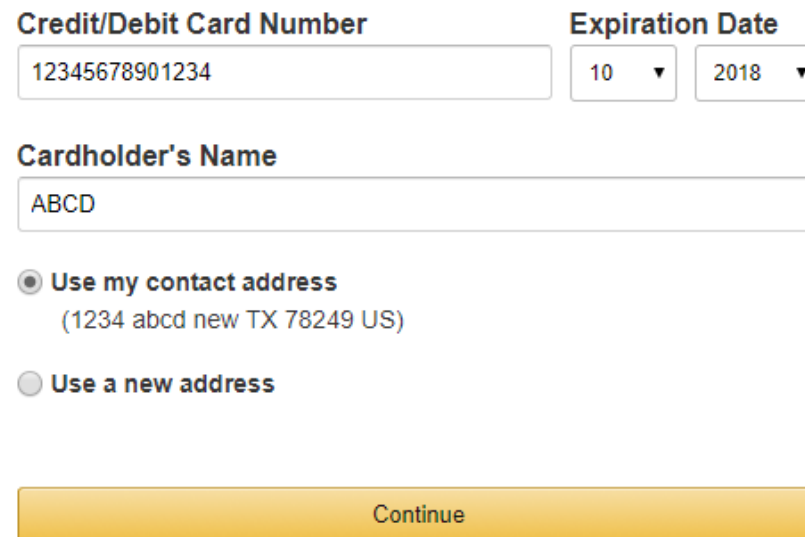
Continue

# Creating an AWS Account (2)

## > Provide required credentials

- >> Enter correct phone number as it will be used to verify credentials
- >> You will be required to enter credit card related information
- >> Your identity will be verified using a telephone number

## > If everything is successful then you will be assigned a 12-digit account number



Credit/Debit Card Number  
12345678901234

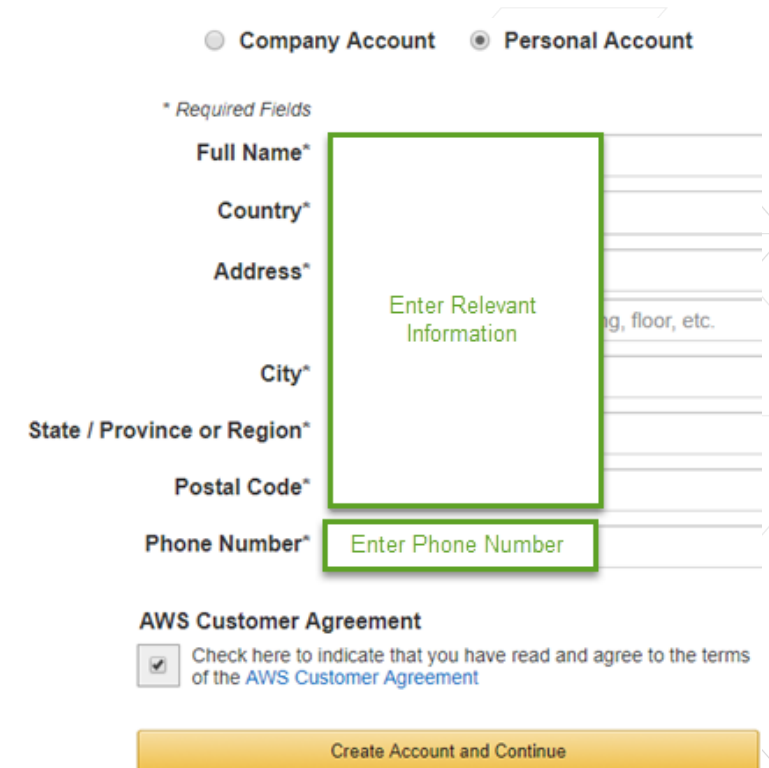
Expiration Date  
10 ▼ 2018 ▼

Cardholder's Name  
ABCD

☒ Use my contact address  
(1234 abcd new TX 78249 US)

☐ Use a new address

Continue



☐ Company Account ☒ Personal Account

\* Required Fields

Full Name\*  
Country\*  
Address\*  
City\*  
State / Province or Region\*  
Postal Code\*  
Phone Number\*

Enter Relevant Information

Enter Phone Number

AWS Customer Agreement  
☒ Check here to indicate that you have read and agree to the terms of the [AWS Customer Agreement](#)

Create Account and Continue

# Signing Up for the Services

- > Open <https://aws.amazon.com/education/awseducate/> and then click **Apply for AWS Educate for Educators** link under the **Educators** section
  - >> Non-academic customers may create an account visiting <https://aws.amazon.com/>, and then choose **Create an AWS Account**
- > **Select the appropriate role (Educators or Students)**

# Signing Up for the Services (2)

- > Fill out Institution Name (you may be able to select), Country, etc. fields
- > Enter the AWS account number that you created in the previous step

<div>Institution Name</div> <div><i>Start typing the name of your school and select from the list. If you don't see your school, enter the full name, example: Harvard University</i></div>	<div>Country</div> <div></div>
<div>--None--</div> <div></div>	<div>City</div> <div></div>
<div>Last Name</div> <div></div>	<div>First Name</div> <div></div>
<div>Title</div> <div><i>Please specify your official title at the institution (eg Assistant Professor)</i></div>	<div>Department</div> <div></div>
	<div>Email</div> <div><i>Please provide a valid, current email issued by your institution. Example: your_name@your_school.edu</i></div>
	<div>AWS Account ID</div> <div><i>Your AWS Account ID is a 12-digit number.</i></div>

# Appendix: Accessing an AWS EC2 instance through your own account





# Login Procedure using Your Own Account

- > Open a web browser and enter <https://signin.aws.amazon.com/console>
- > Enter email address of your AWS account or account ID and click Next
- > Enter password and click on Sign-In
- > Select N. Virginia farm using the drop-down button



Sign in ⓘ

Email address of your AWS account

To sign in as an IAM user, enter your account ID or account alias instead.

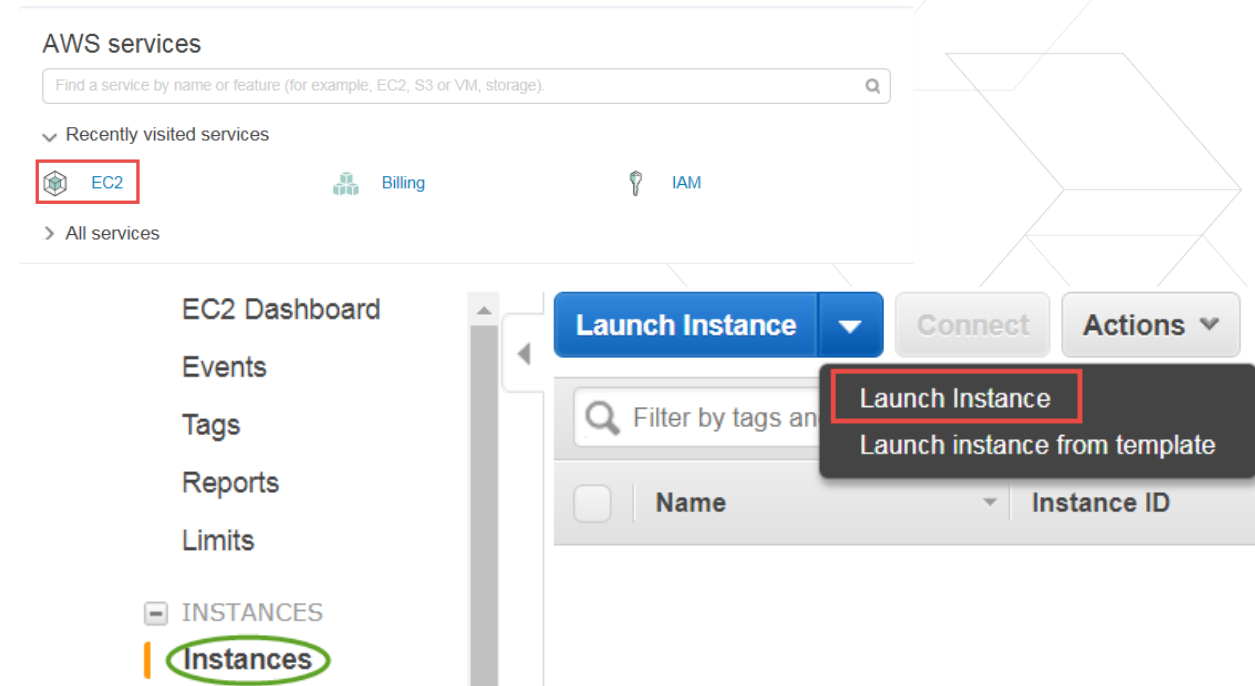
Next

Password

Sign In


# Starting an Instance on EC2

- > You can use an available AMI to start an instance or you can use a provided AMI to start an instance
- > Click on the EC2 icon
  - >> EC2 dashboard will be displayed
- > Click on the Instances option under the INSTANCES group
- > Click on the Launch Instance > Launch Instance
  - >> The seven-steps wizard will open



# Choosing AMI

- > Step 1: Select AWS Marketplace
- > In the search field enter *FPGA* and hit enter

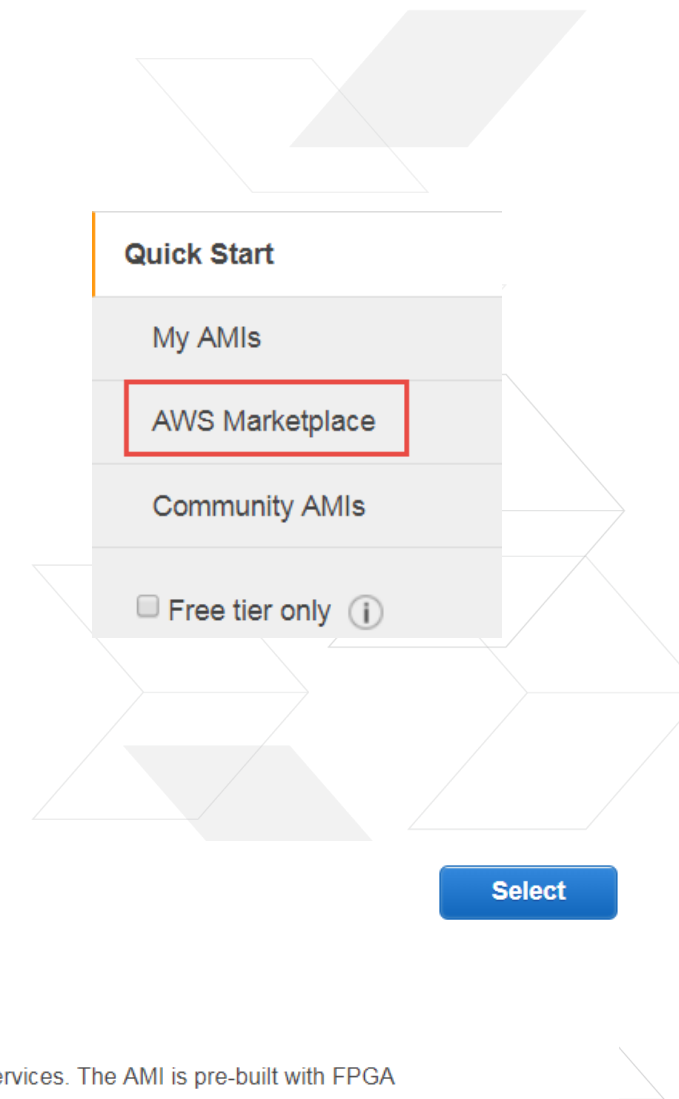


- > Click on the **Select** button of the **FPGA Developer AMI**
  - >> List of instance types with cost be displayed (see next slide)



Free tier eligible

**FPGA Developer AMI**  
★★★★★ (3) | 1.3.3 | Sold by [Amazon Web Services](#)  
\$0.00/hr for software + AWS usage fees  
Linux/Unix, CentOS 7.3 | 64-bit Amazon Machine Image (AMI) | Updated: 9/24/17  
The FPGA (field programmable gate array) AMI is a supported and maintained CentOS Linux image provided by Amazon Web Services. The AMI is pre-built with FPGA development tools and ...  
[More info](#)



## Pricing Details

Pricing may vary

### Hourly Fees

Instance Type	Software	EC2	Total
T2 Nano	\$0.00	\$0.006	<b>\$0.006/hr</b>
R4 16 Extra Large	\$0.00	\$4.256	<b>\$4.256/hr</b>
M5 Extra Large	\$0.00	\$0.192	<b>\$0.192/hr</b>
M4 Extra Large	\$0.00	\$0.20	<b>\$0.20/hr</b>
H1 2 Extra Large	\$0.00	\$0.55	<b>\$0.55/hr</b>
High I/O Quadruple Extra Large	\$0.00	\$1.248	<b>\$1.248/hr</b>
T2 Large	\$0.00	\$0.093	<b>\$0.093/hr</b>
C4 Double Extra Large	\$0.00	\$0.398	<b>\$0.398/hr</b>
M5 Large	\$0.00	\$0.096	<b>\$0.096/hr</b>
FPGA Accelerated Compute 16 Extra Large	\$0.00	\$13.20	<b>\$13.20/hr</b>
C5 Large	\$0.00	\$0.085	<b>\$0.085/hr</b>
M5 Double Extra Large	\$0.00	\$0.384	<b>\$0.384/hr</b>
T2 Double Extra Large	\$0.00	\$0.371	<b>\$0.371/hr</b>
T2 Extra Large	\$0.00	\$0.186	<b>\$0.186/hr</b>
C4 Eight Extra Large	\$0.00	\$1.591	<b>\$1.591/hr</b>
M4 Quadruple Extra Large	\$0.00	\$0.80	<b>\$0.80/hr</b>
M5 12 Extra Large	\$0.00	\$2.304	<b>\$2.304/hr</b>
T2 Medium	\$0.00	\$0.046	<b>\$0.046/hr</b>
C4 Large	\$0.00	\$0.10	<b>\$0.10/hr</b>

T2 Small	\$0.00	\$0.023	<b>\$0.023/hr</b>
R4 Extra Large	\$0.00	\$0.266	<b>\$0.266/hr</b>
C4 Quadruple Extra Large	\$0.00	\$0.796	<b>\$0.796/hr</b>
C4 Extra Large	\$0.00	\$0.199	<b>\$0.199/hr</b>
R4 Eight Extra Large	\$0.00	\$2.128	<b>\$2.128/hr</b>
C5 Double Extra Large	\$0.00	\$0.34	<b>\$0.34/hr</b>
R4 Double Extra Large	\$0.00	\$0.532	<b>\$0.532/hr</b>
High I/O Double Extra Large	\$0.00	\$0.624	<b>\$0.624/hr</b>
D2 Extra Large	\$0.00	\$0.69	<b>\$0.69/hr</b>
D2 Eight Extra Large	\$0.00	\$5.52	<b>\$5.52/hr</b>
T2 Micro	\$0.00	\$0.012	<b>\$0.012/hr</b>
C5 Extra Large	\$0.00	\$0.17	<b>\$0.17/hr</b>
C5 Eighteen Extra Large	\$0.00	\$3.06	<b>\$3.06/hr</b>
C5 Nine Extra Large	\$0.00	\$1.53	<b>\$1.53/hr</b>
R4 Quadruple Extra Large	\$0.00	\$1.064	<b>\$1.064/hr</b>
D2 Double Extra Large	\$0.00	\$1.38	<b>\$1.38/hr</b>
M4 Large	\$0.00	\$0.10	<b>\$0.10/hr</b>
FPGA Accelerated Compute Double Extra Large	\$0.00	\$1.65	<b>\$1.65/hr</b>
M5 24 Extra Large	\$0.00	\$4.608	<b>\$4.608/hr</b>
H1 4 Extra Large	\$0.00	\$1.10	<b>\$1.10/hr</b>

Red-boxed are FPGA based

Green-boxed are compute optimized

Perform development on compute-optimized instances, perform final verification in hardware on FPGA based instances

# Instance Configuration

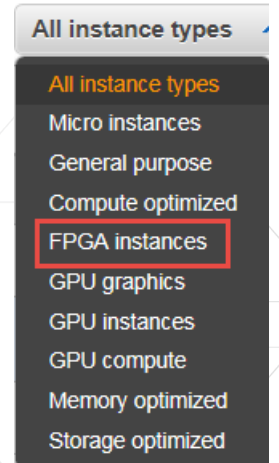
- > In Step 2. Click *All instance type* filter and select FPGA instances
  - >> Two entries will be displayed
- > Select c4.4xlarge or f1.2xlarge type and click **Next: Configure Instance Details**

Filter by: **FPGA instances** **Current generation** [Show/Hide Columns](#)

Currently selected: f1.2xlarge (26 ECUs, 8 vCPUs, 2.3 GHz, Intel Xeon E5-2686v4, 122 GiB memory, 1 x 470 GiB Storage Capacity)

	Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
<input checked="" type="checkbox"/>	FPGA instances	f1.2xlarge	8	122	1 x 470 (SSD)	Yes	Up to 10 Gigabit
<input type="checkbox"/>	FPGA instances	f1.16xlarge	64	976	4 x 940 (SSD)	Yes	25 Gigabit

[Cancel](#) [Previous](#) [Review and Launch](#) **Next: Configure Instance Details**



# Instance Configuration (2)

> Step 3: *Configure Instance Details* will be displayed

> Click Next: Add Storage

> Step 4: *Add Storage* will be displayed

> Click Next: Add Tags

> Step 5: *Add Tags* will be displayed

> Click Next: Configure Security Group

> Step 6: *Configure Security Group* will be displayed

- >> Select *Create a new security group* if this is first time you are creating an instance on AWS
- >> Select *Select an existing security group* if you already have created earlier

The screenshot displays the 'Step 6: Configure Security Group' interface in the AWS Management Console. At the top, there is a table of security rules with columns: Type, Protocol, Port Range, Source, and Description. The rules listed are HTTP (TCP, 80), Custom TCP Rule (TCP, 8080), SSH (TCP, 22), and RDP (TCP, 3389). Below the table are buttons for 'Cancel', 'Previous', and 'Review and Launch'. The main text explains that a security group is a set of firewall rules and provides instructions on how to add rules. Under the 'Assign a security group' section, there are two radio buttons: 'Create a new security group' and 'Select an existing security group'. The 'Select an existing security group' option is chosen. Below this, a table lists existing security groups with columns: Security, Name, and Description. The security group 'sg-68570c1a ICCD\_Tutorial' is selected and highlighted with a red box.

Type	Protocol	Port Range	Source	Description
HTTP	TCP	80	0.0.0.0/0	
Custom TCP Rule	TCP	8080	0.0.0.0/0	
SSH	TCP	22	0.0.0.0/0	
RDP	TCP	3389	0.0.0.0/0	

Step 6: Configure Security Group

A security group is a set of firewall rules that control the traffic for your instance. On this page, you can add rules to allow specific traffic to reach your instance. For example, if you want to set up a web server and allow Internet traffic to reach your instance, add rules that allow unrestricted access to the HTTP and HTTPS ports. You can create a new security group or select from an existing one below. [Learn more](#) about Amazon EC2 security groups.

Assign a security group: ☐ Create a new security group ☒ Select an existing security group

Security	Name	Description
<input type="radio"/>	sg-9798a3e9default	default VPC security group
<input type="radio"/>	sg-f8473a8a FPGA Developer AMI-1-3-3-AutogenByAWSMP-3	This security group was generated by AWS Marketplace
<input checked="" type="radio"/>	sg-68570c1a ICCD_Tutorial	For ICCD Tutorial

# Instance Configuration (3)

## > Click on Add Rule and then add desired ports and type of connections

>> Add 3389 port for the RDP (Remote Desktop)

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ	Description ⓘ
HTTP	TCP	80	0.0.0.0/0	
Custom TCP Rule	TCP	8080	0.0.0.0/0	
SSH	TCP	22	0.0.0.0/0	
RDP	TCP	3389	0.0.0.0/0	

Cancel Previous **Review and Launch**

## > Click on Review and Launch

## > Step 7: *Review Instance Launch* will be displayed

>> This summarizes the AMI details, Instance type, Security Group etc.

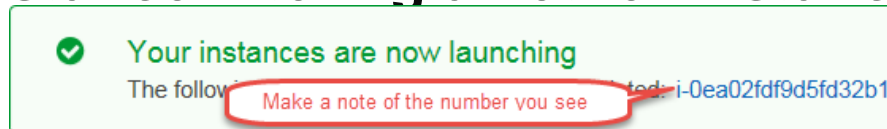
>> Ignore warnings of non-free usage tier and security group

## > Click Launch

>> A form will be displayed to either create or select a key pair

# Instance Configuration (4)

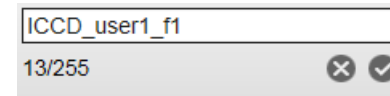
- > **Click the acknowledgement check box, and click Launch Instances**
  - >> A Launch Status be displayed
- > **Click on the link and it will take you to the instance filtering all other instances**





# Naming the Instance

- > You can name the instance by moving the mouse under the Name column until you see a pencil tool
- > Click on the pencil tool and a field will open
- > Enter a meaningful name and hit enter

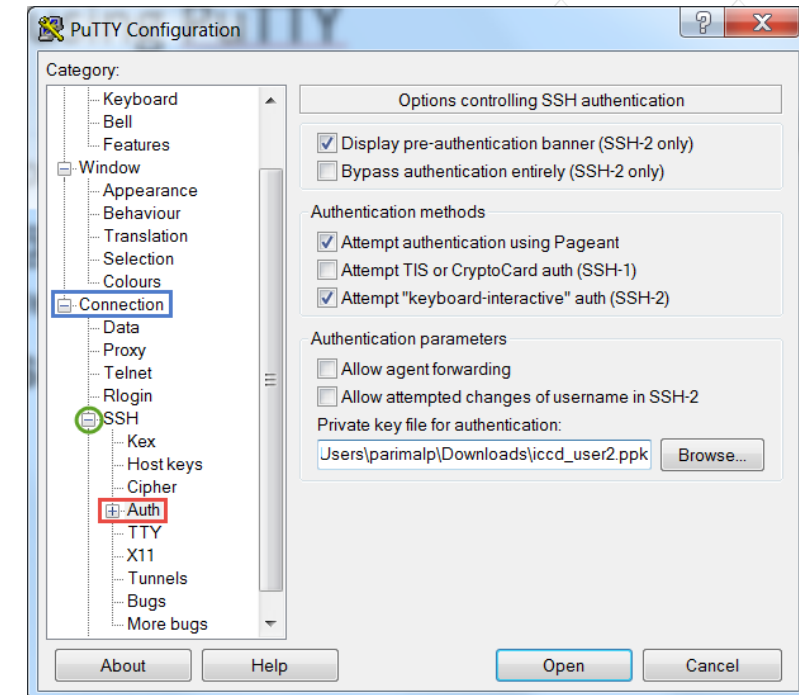
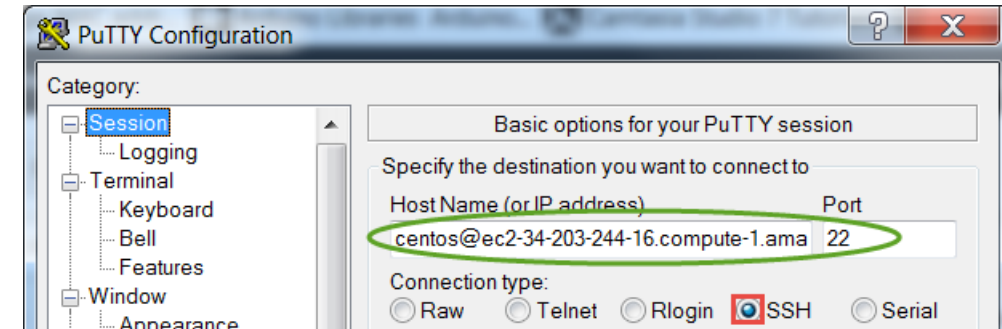


# Appendix: Other Ways to Communicate



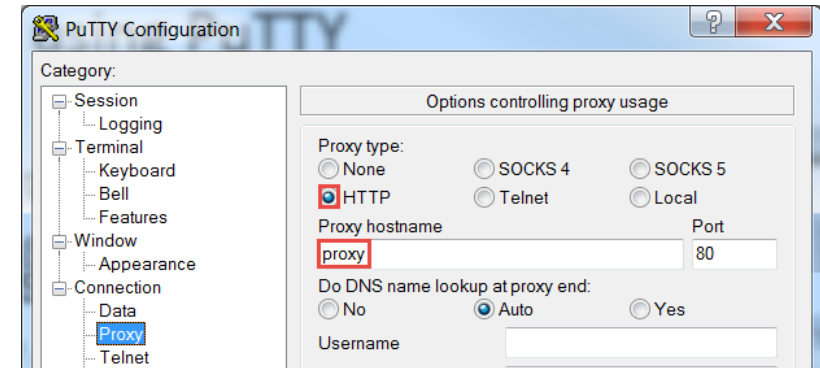
# Connecting using PuTTY (Windows)

- > Start PuTTY program
- > Enter `centos@<public_dns_entry>` in the *Host Name* field and 22 in the *Port* field
- > Make sure SSH is selected as the *Connection Type*
- > Expand *SSH* under the *Connection* and click *Auth*
- > Click on the Browse button and browse to the location where you were provided the private key
- > Click Open
- > Click Yes



# Connecting using PuTTY on VPN (Windows)

- > Start PuTTY program
- > Enter `centos@<public_dns_entry>` in the *Host Name* field and 22 in the *Port* field
- > Make sure SSH is selected as the *Connection Type*
- > Select Proxy under the Connection
- > Select HTTP as the *Proxy type*, proxy as the *Proxy hostname* and 80 as the *Port* number
- > Expand *SSH* under the *Connection* and click Auth
- > Click on the Browse button and browse to the location of the private key
- > Click Open and then click Yes



## Connecting using Terminal (MAC)

- > **Start a Terminal program**
  - >> Click on Finder, select Applications on left
  - >> Enter Terminal in the Search bar and double-click on the entry
- > **Make sure that the access mode to the \*.pem file is 400**
- > **Enter the ssh command**  

```
ssh -i <pem file> centos@<public_dns ent
```
- > **At the centos prompt set the password if not already set**
  - >> This is needed to enable remote desktop connection: `sudo passwd centos`
  - >> Enter user password as the centos password
- > **At this stage you can execute command line tasks**

[illegible]

**Adaptable.**  
**Intelligent.**

